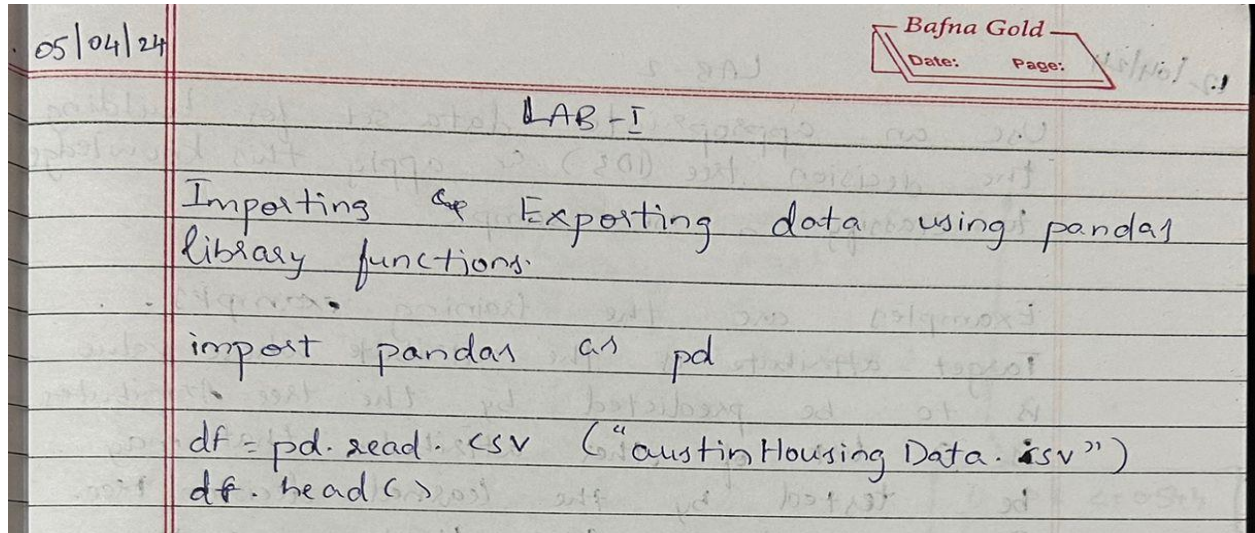


Lab1

Date: 5th April , 2024

Observation

Program Title: 1. Write a python program to import and export data using Pandas library functions



Program Title: 1. Write a python program to import and export data using Pandas library Functions

```
# Read data from URL
iris_data = pd.read_csv(url, names=col_names)

iris_data.head()

# Export the file to the current working directory
iris_data.to_csv("cleaned_iris_data.csv")
```

Program Title: 2. Demonstrate various data pre-processing techniques for a given dataset

Code

```
# import the pandas library
import pandas as pd

# Read the CSV file
airbnb_data = pd.read_csv("/content/sample_data/mnist_test.csv")

# View the first 5 rows
airbnb_data.head()

# Webpage URL
url =
"https://archive.ics.uci.edu/ml/machine-learning-databases/iris/iris.data"

# Define the column names
col_names = ["sepal_length_in_cm",
             "sepal_width_in_cm",
             "petal_length_in_cm",
             "petal_width_in_cm",
             "class"]

# Read data from URL
iris_data = pd.read_csv(url, names=col_names)

iris_data.head()

# Export the file to the current working directory
iris_data.to_csv("cleaned_iris_data.csv")
```

Program Title: 2. Demonstrate various data pre-processing techniques for a given dataset

* Reading data from URL

```
url = "https://arch.ur.ics.vci.edu/me/  
machinelearning/databases/iris/iris.data"
```

```
col_names = ["sepal-length-in-cm", "sepal-  
width-in-cm", "petal-length-in-cm", "petal-  
width-in-cm", "class?"]
```

```
iris_data = pd.read_csv(url, names=col_names)  
iris_data.head()
```

* Exporting to another csv file

```
iris_data.to_csv("cleaned-iris-  
data.csv")
```

Snapshot of the output

```
[5] import pandas as pd

# Read the CSV file
airbnb_data = pd.read_csv("/content/sample_data/mnist_test.csv")

# View the first 5 rows
airbnb_data.head()
```

	7	0	0.1	0.2	0.3	0.4	0.5	0.6	0.7	0.8	...	0.658	0.659	0.660	0.661	0.662	0.663	0.664	0.665	0.666	0.667
0	2	0	0	0	0	0	0	0	0	0	...	0	0	0	0	0	0	0	0	0	0
1	1	0	0	0	0	0	0	0	0	0	...	0	0	0	0	0	0	0	0	0	0
2	0	0	0	0	0	0	0	0	0	0	...	0	0	0	0	0	0	0	0	0	0
3	4	0	0	0	0	0	0	0	0	0	...	0	0	0	0	0	0	0	0	0	0
4	1	0	0	0	0	0	0	0	0	0	...	0	0	0	0	0	0	0	0	0	0

5 rows x 785 columns

```
# Webpage URL
url = "https://archive.ics.uci.edu/ml/machine-learning-databases/iris/iris.data"

# Define the column names
col_names = ["sepal_length_in_cm",
             "sepal_width_in_cm",
             "petal_length_in_cm",
             "petal_width_in_cm",
             "class"]

# Read data from URL
iris_data = pd.read_csv(url, names=col_names)

iris_data.head()
```

	sepal_length_in_cm	sepal_width_in_cm	petal_length_in_cm	petal_width_in_cm	class
0	5.1	3.5	1.4	0.2	Iris-setosa
1	4.9	3.0	1.4	0.2	Iris-setosa
2	4.7	3.2	1.3	0.2	Iris-setosa
3	4.6	3.1	1.5	0.2	Iris-setosa
4	5.0	3.6	1.4	0.2	Iris-setosa

Next steps: [View recommended plots](#)

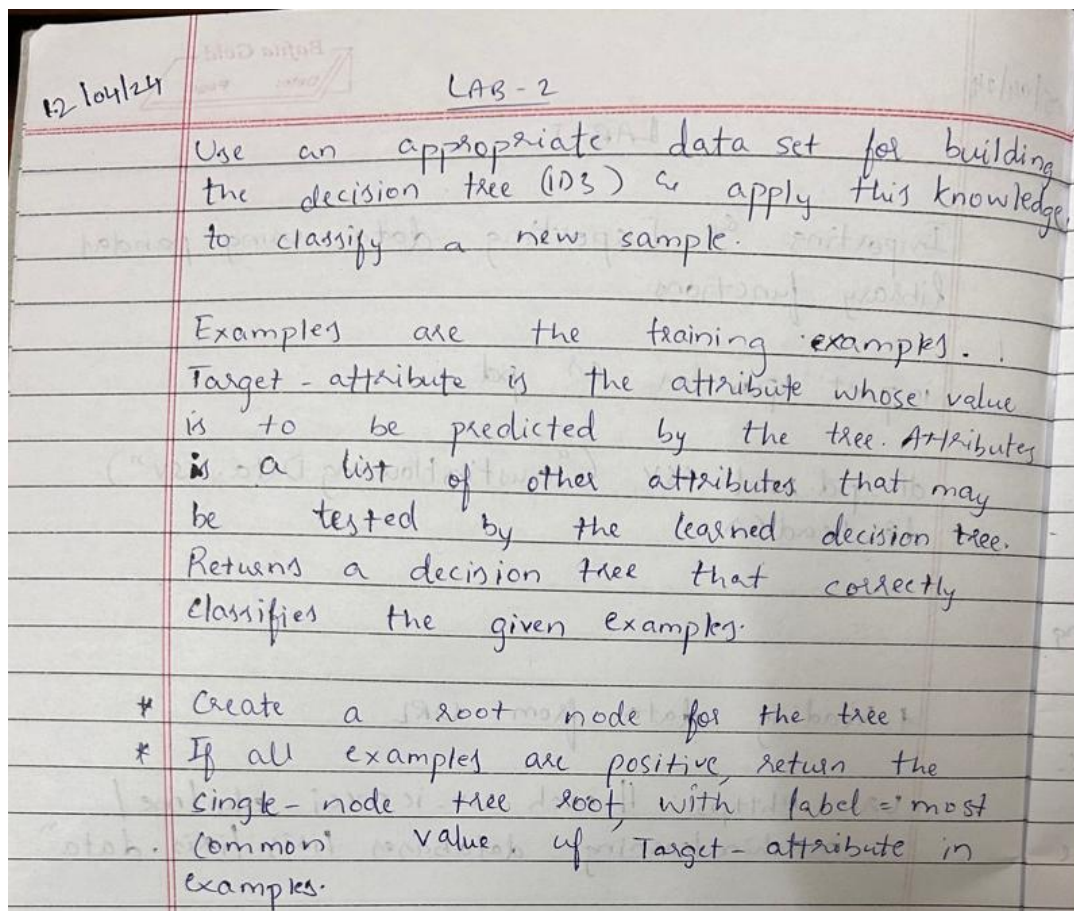
```
[8] # Export the file to the current working directory
iris_data.to_csv("cleaned_iris_data.csv")
```

Lab2

Date: 12th April, 2024

Title: Use an appropriate data set for building the decision tree (ID3) and apply this knowledge to classify a new sample.

Algorithm



Code and output

```
import the pandas library
import pandas as pd
```

```
# Read the CSV file
airbnb_data = pd.read_csv("../content/sample_data/mnist_test.csv")

# View the first 5 rows
airbnb_data.head()
```

	7	0	0.1	0.2	0.3	0.4	0.5	0.6	0.7	0.8	...	0.658	0.659	0.660	0.661	0.662	0.663	0.664	0.665	0.666	0.667
0	2	0	0	0	0	0	0	0	0	0	...	0	0	0	0	0	0	0	0	0	0
1	1	0	0	0	0	0	0	0	0	0	...	0	0	0	0	0	0	0	0	0	0
2	0	0	0	0	0	0	0	0	0	0	...	0	0	0	0	0	0	0	0	0	0
3	4	0	0	0	0	0	0	0	0	0	...	0	0	0	0	0	0	0	0	0	0
4	1	0	0	0	0	0	0	0	0	0	...	0	0	0	0	0	0	0	0	0	0

5 rows × 785 columns

```
# Webpage URL
url = "https://archive.ics.uci.edu/ml/machine-learning-databases/iris/iris.data"

# Define the column names
col_names = ["sepal_length_in_cm",
             "sepal_width_in_cm",
             "petal_length_in_cm",
             "petal_width_in_cm",
             "class"]

# Read data from URL
iris_data = pd.read_csv(url, names=col_names)

iris_data.head()
```

	sepal_length_in_cm	sepal_width_in_cm	petal_length_in_cm	petal_width_in_cm	class
0	5.1	3.5	1.4	0.2	Iris-setosa
1	4.9	3.0	1.4	0.2	Iris-setosa
2	4.7	3.2	1.3	0.2	Iris-setosa
3	4.6	3.1	1.5	0.2	Iris-setosa
4	5.0	3.6	1.4	0.2	Iris-setosa

```
iris_data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 150 entries, 0 to 149
Data columns (total 5 columns):
#   Column              Non-Null Count  Dtype
---  ---
0   sepal_length_in_cm    150 non-null    float64
1   sepal_width_in_cm     150 non-null    float64
2   petal_length_in_cm    150 non-null    float64
3   petal_width_in_cm     150 non-null    float64
4   class                 150 non-null    object
dtypes: float64(4), object(1)
memory usage: 6.0+ KB
```

```
iris_data.describe()
```

	sepal_length_in_cm	sepal_width_in_cm	petal_length_in_cm	petal_width_in_cm
count	150.000000	150.000000	150.000000	150.000000
mean	5.843333	3.054000	3.758667	1.198667
std	0.828066	0.433594	1.764420	0.763161
min	4.300000	2.000000	1.000000	0.100000
25%	5.100000	2.800000	1.600000	0.300000
50%	5.800000	3.000000	4.350000	1.300000
75%	6.400000	3.300000	5.100000	1.800000
max	7.900000	4.400000	6.900000	2.500000

```
iris_data.isnull().sum()
```

sepal_length_in_cm	0
sepal_width_in_cm	0
petal_length_in_cm	0
petal_width_in_cm	0
class	0

dtype: int64

```
data=iris_data.to_numpy()
dataset=data[:, :-1]
df = pd.DataFrame(dataset, index=dataset[:, 0])
df.kurt(axis=1)
```

5.1	-2.368842
4.9	-1.891924
4.7	-2.276657
4.6	-1.57517
5.0	-2.787804
...	
6.7	-2.983686
6.3	-3.798183
6.5	-3.127297
6.2	-3.387994
5.9	-3.345923

Length: 150, dtype: object

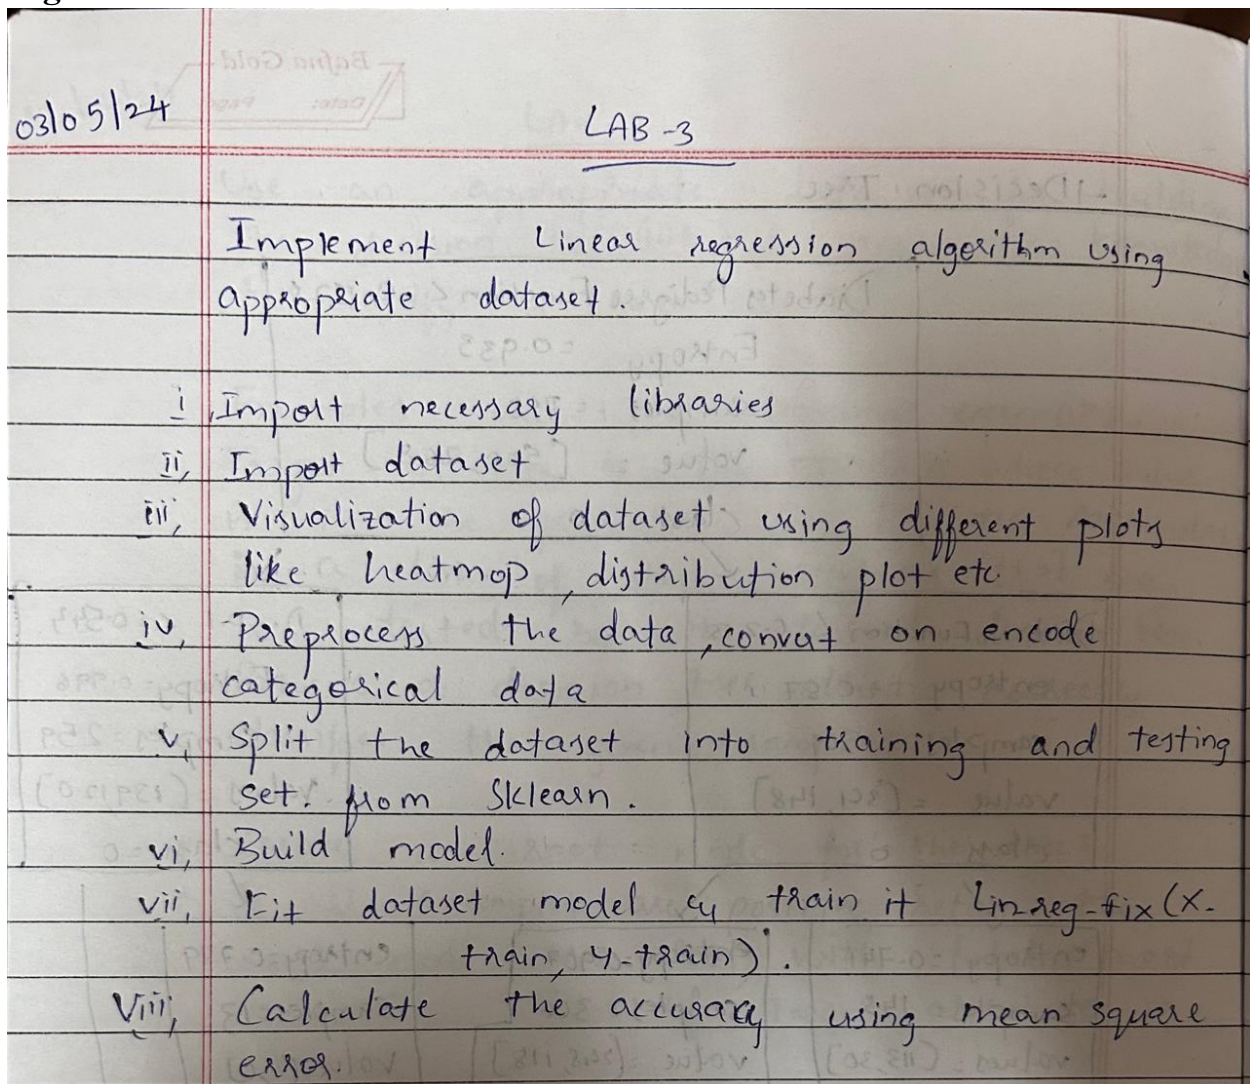
```
# Export the file to the current working directory
iris_data.to_csv("cleaned_iris_data.csv")
```

Lab3

Date: 3rd May, 2024

Title: Implement Linear Regression algorithm using appropriate dataset

Algorithm



Code and output


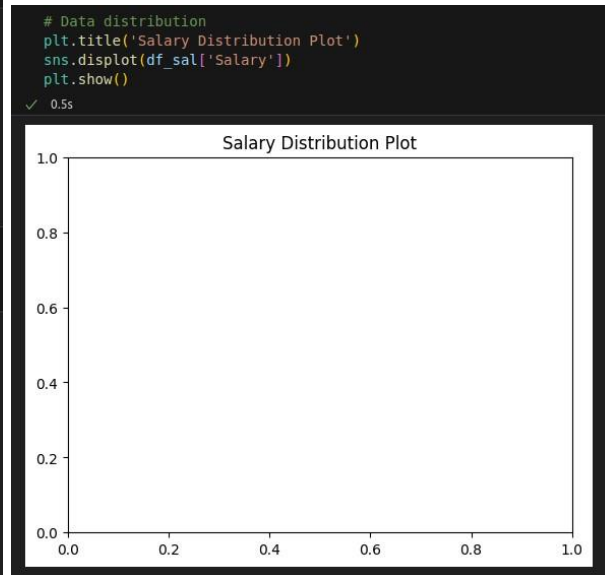
```
# Import libraries (module) pd
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from pandas.core.common import random_state
from sklearn.linear_model import LinearRegression

# Get dataset
df_sal = pd.read_csv('salary_data.csv')
df_sal.head()
```

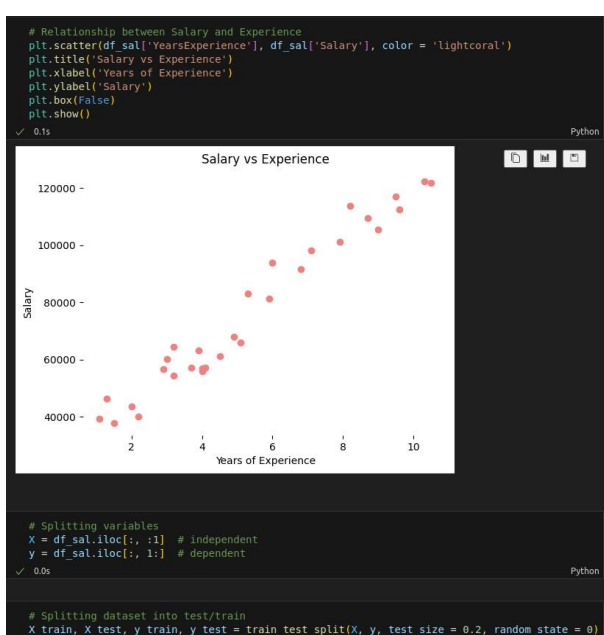
	YearsExperience	Salary
0	1.1	39343
1	1.3	46205
2	1.5	37731
3	2.0	43525
4	2.2	39891

```
# Describe data
df_sal.describe()
```

	YearsExperience	Salary
count	30.000000	30.000000
mean	5.313333	76003.000000
std	2.837888	27414.429785
min	1.100000	37731.000000
25%	3.200000	56720.750000
50%	4.700000	65237.000000
75%	7.700000	100544.750000
max	10.500000	122391.000000



```
# Relationship between Salary and Experience
plt.scatter(df_sal['YearsExperience'], df_sal['Salary'], color = 'lightcoral')
plt.title('Salary vs Experience')
plt.xlabel('Years of Experience')
plt.ylabel('Salary')
plt.box(False)
plt.show()
```




```
# Regressor model
regressor = LinearRegression()
regressor.fit(X_train, y_train)
```

✓ 0.2s

Python

```
LinearRegression()
```

```
# Prediction result
y_pred_test = regressor.predict(X_test) # predicted value of y_test
y_pred_train = regressor.predict(X_train) # predicted value of y_train
```

✓ 0.0s


Python

```
# Prediction on training set
plt.scatter(X_train, y_train, color = 'lightcoral')
plt.plot(X_train, y_pred_train, color = 'firebrick')
plt.title('Salary vs Experience (Training Set)')
plt.xlabel('Years of Experience')
plt.ylabel('Salary')
plt.legend(['X_train/Pred(y_test)', 'X_train/y_train'], title = 'Sal/Exp', loc='best', facecolor='white')
plt.box(False)
plt.show()
```



```
# Prediction on test set
plt.scatter(X_test, y_test, color = 'lightcoral')
plt.plot(X_train, y_pred_train, color = 'firebrick')
plt.title('Salary vs Experience (Test Set)')
plt.xlabel('Years of Experience')
plt.ylabel('Salary')
plt.legend(['X_train/Pred(y_test)', 'X_train/y_train'], title = 'Sal/Exp', loc='best', facecolor='white')
plt.box(False)
plt.show()
```

✓ 0.0s



```
# Regressor coefficients and intercept
print(f'Coefficient: {regressor.coef_}')
print(f'Intercept: {regressor.intercept_}')
```

✓ 0.0s

```
Coefficient: [[9312.57512673]]
Intercept: [26788.09915063]
```

Title: Implement Multi-Linear Regression algorithm using appropriate dataset

Algorithm

b) Implement Multi-Linear Regression.

Algorithm:

- 1) Import the required python package
- 2) Load the dataset
- 3) Do the data analysis
- 4) Split the dataset into dependent / independent variables.
- 5) One-hot encoding of categorical data.
* It is a method to represent a categorical variable in a numerical way.
- 6) Split data into train/test sets
- 7) Train the regression model
- 8) Predict the results

Code and output

```
#Importing the libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

# import warnings
import warnings
warnings.filterwarnings("ignore")

# We will use some methods from the sklearn module
from sklearn import linear_model
from sklearn.linear_model import LinearRegression
from sklearn import metrics
from sklearn.metrics import mean_squared_error, mean_absolute_error
from sklearn.model_selection import train_test_split, cross_val_score
```

```
[ ] # Reading the Dataset
df = pd.read_csv("data.csv")
```

```
[ ] df.head()
```

```
[ ]
```

	Car	Model	Volume	Weight	CO2
0	Toyota	Aygo	1000	790	99
1	Mitsubishi	Space Star	1200	1100	95
2	Skoda	Citigo	1000	929	95
3	Fiat	500	900	865	90
4	Mini	Cooper	1500	1140	105

```
[ ] df.shape
```

```
[ ] (36, 5)
```

```
[ ] df.corr(numeric_only=True)
```

```
[ ]
```

	Volume	Weight	CO2
Volume	1.000000	0.753537	0.592082
Weight	0.753537	1.000000	0.552150
CO2	0.592082	0.552150	1.000000


```
[ ] print(df.describe())
```

```
[ ]
```


	Volume	Weight	CO2
count	36.000000	36.000000	36.000000
mean	1611.111111	1292.277778	102.627778
std	388.975047	242.123889	7.454571
min	900.000000	790.000000	90.000000
25%	1475.000000	1117.250000	97.750000
50%	1600.000000	1329.000000	99.000000
75%	2000.000000	1418.250000	105.000000
max	2500.000000	1746.000000	120.000000

```
[ ] #Setting the value for X and Y
X = df[['Weight', 'Volume']]
y = df['CO2']
```


```
[ ] fig, axes = plt.subplots(2, figsize = (5,5))
plt1 = sns.boxplot(df['Weight'], ax = axes[0])
plt2 = sns.boxplot(df['Volume'], ax = axes[1])
plt.tight_layout()
```




```
[ ] sns.distplot(df['CO2']);
```



```
[ ] sns.pairplot(df, x_vars=['Weight', 'Volume'], y_vars='CO2', height=4, aspect=1, kind='scatter')
plt.show()
```



```
[ ] # Create the correlation matrix and represent it as a heatmap.
sns.heatmap(df.corr(numeric_only=True), annot = True, cmap = 'coolwarm')
plt.show()
```



```
[ ] X_train,X_test, y_train, y_test = train_test_split(X, y, test_size = 0.3, random_state = 100)
```

```
[ ] y_train.shape
```

```
(25,)
```

```
[ ] y_test.shape
```

```
(11,)
```

```
reg_model = linear_model.LinearRegression()
```

```
[ ] #Fitting the Multiple Linear Regression model
reg_model = LinearRegression().fit(X_train, y_train)
```

```
[ ] #Printing the model coefficients
print('Intercept: ',reg_model.intercept_)
# pair the feature names with the coefficients
list(zip(X, reg_model.coef_))
```

```
Intercept: 74.33882836589245
[('Weight', 0.0171800645996374), ('Volume', 0.0025846399866482976)]
```

```
[ ] #Predicting the Test and Train set result
y_pred= reg_model.predict(X_test)
x_pred= reg_model.predict(X_train)
```

```
[ ] print("Prediction for test set: {}".format(y_pred))
```

```
Prediction for test set: [ 90.41571939 102.16323413 99.56363213 104.56661845 101.54657652
 95.94778019 108.64011848 102.22654214 92.80374837 97.27327129
 97.57074463]
```

```
[ ] #Actual value and the predicted value
reg_model_diff = pd.DataFrame({'Actual value': y_test, 'Predicted value': y_pred})
reg_model_diff
```

	Actual value	Predicted value
0	99	90.415719
19	105	102.163234
32	104	99.563632
35	120	104.566618
7	92	101.546577
12	99	95.947700
29	114	108.640118
33	108	102.226542
5	105	92.803748
1	95	97.273271
18	104	97.570745

```
mae = metrics.mean_absolute_error(y_test, y_pred)
mse = metrics.mean_squared_error(y_test, y_pred)
r2 = np.sqrt(metrics.mean_squared_error(y_test, y_pred))

print('Mean Absolute Error:', mae)
print('Mean Square Error:', mse)
print('Root Mean Square Error:', r2)
```

```
Mean Absolute Error: 6.901980901636316
Mean Square Error: 63.39765310998794
Root Mean Square Error: 7.96226432053018
```


Title: Build KNN Classification model for a given dataset

Algorithm

c) Implementation of KNN - Algorithm.

- 1) ^{Import} ~~Importing~~ the modules
- 2) Creating Dataset
- 3) Visualise the Dataset
- 4) Splitting data into training & testing datasets
- 5) KNN classify implementation
- 6) Prediction the KNN classifiers
- 7) Predict Accuracy for both k values

OUTPUT:

Accuracy with $k=5 \Rightarrow 93.600001$

Accuracy with $k=1 \Rightarrow 90.4$

Code and output

```
from sklearn.model_selection import train_test_split
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import classification_report, confusion_matrix
from sklearn import datasets
iris = datasets.load_iris()

x = iris.data
y = iris.target

print('sepal-length', 'sepal-width', 'petal-length', 'petal-width')
print(x)
print('class: 0 - Iris-Setosa, 1 - Iris-Versicolour, 2 - Iris-Virginica')
print(y)
```

[illegible]

```
x_train, x_test, y_train, y_test = train_test_split(x,y, test_size=0.3)

#To Training the model and Nearest neighbors K=5
classifier = KNeighborsClassifier(n_neighbors=5)
classifier.fit(x_train, y_train)

#to make predictions on our test data
y_pred=classifier.predict(x_test)

print('Prediction - ')

for i, test in enumerate(x_test) :
    print(f'{i} {test} - {y_pred[i]}')

# print('Confusion Matrix')
# print(confusion_matrix(y_test,y_pred))
# print('Accuracy Metrics')
# print(classification_report(y_test,y_pred))
```

```

[5.2 4.1 1.1 5.0 1.] - 0
[5.5 2.3 4. 1.3] -1
[6.7 3.1 4.7 1.5] -1
[7. 3.2 4.7 1.4] -1
[6.2 2.8 4.8 1.8] -2
[5.7 2.8 4.5 1.3] -2
[6. 3.4 4.5 1.6] -1
[5.1 3.8 1.6 1.02] -0
[5.5 2.5 4. 1.3] -1
[4.8 3.1 1.6 1.02] -0
[6.1 3. 4.9 1.8] -2
[4.7 3.2 1.6 1.02] -0
[5.6 2.9 3.6 1.3] -1
[5.4 3.9 1.3 0.4] -0
[5.3 3.2 1.2 0.2] -0
[6.1 2.9 4.7 1.4] -1
[5.6 3.2 1.6 1.02] -0
[6.7 2.8 6.7 2. 2] -2
[4.6 3.2 1.4 1.02] -0
[5.7 2.9 4.2 1.3] -1
[4.6 3.6 1. 1.02] -0
[6.8 2.8 4.8 1.4] -1
[6.8 3.2 5.9 1.3] -2
[6.2 2.2 4.5 1.5] -1
...
[5.2 3.4 1.4 1.02] -0
[6.1 2.8 4.7 1.2] -1
[6.5 3. 5.2 2.] -2
[6.5 3. 5.8 2.2] -2

```

Lab4

Date: 17th May , 2024

Title: Build Logistic Regression Model for a given dataset

Algorithm

03/05/24

LAB-4

Bafna Gold
Date: Page:

1. Build Logistic Regression Model

- 1) Import required libraries
- 2) Load, visualize and explore the dataset
- 3) Clean the dataset
- 4) Deal with the outliers
- 5) Define dependent & independent variables and then split the data into a training set and testing set.

OUTPUT

Regression coefficients obtained are $b_0 = -68.83$
 $b_1 = 0.192671$

Code and output

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from plotly.offline import download_plotlyjs, init_notebook_mode, plot, iplot
import plotly as py
import plotly.graph_objs as go
import time

init_notebook_mode(connected=True)
```

```
In [2]: def sigmoid(X, weight):
z = np.dot(X, weight)
return 1 / (1 + np.exp(-z))
```

```
In [3]: def loss(h, y):
return (-y * np.log(h) - (1 - y) * np.log(1 - h)).mean()
```

```
In [4]: def gradient_descent(X, h, y):
return np.dot(X.T, (h - y)) / y.shape[0]
def update_weight_loss(weight, learning_rate, gradient):
return weight - learning_rate * gradient
```

```
In [5]: def log_likelihood(x, y, weights):
z = np.dot(x, weights)
ll = np.sum(y*z - np.log(1 + np.exp(z)))
return ll
```

```
In [6]: def gradient_ascent(X, h, y):
return np.dot(X.T, y - h)
def update_weight_ascent(weight, learning_rate, gradient):
return weight + learning_rate * gradient
```

```
In [7]: data = pd.read_csv('/content/WA_Fn-UseC-Telco-Customer-Churn.csv')
print("Dataset size")
print("Rows {} Columns {}".format(data.shape[0], data.shape[1]))
print("Columns and data types")
pd.DataFrame(data.dtypes).rename(columns = {'0': 'dtype'})
```

Dataset size
Rows 7043 Columns 21
Columns and data types

```
Out[7]:
```

	dtype
customerID	object
gender	object
SeniorCitizen	int04
Partner	object
Dependents	object
tenure	int04
PhoneService	object
MultipleLines	object
InternetService	object
OnlineSecurity	object
OnlineBackup	object
DeviceProtection	object
TechSupport	object
StreamingTV	object
StreamingMovies	object
Contract	object
PaperlessBilling	object
PaymentMethod	object
MonthlyCharges	float04
TotalCharges	object
Churn	object

```
In [8]: df = data.copy()
```

```
In [14]: churns = ["Yes", "No"]
fig = {
    'data': [
        {
            'x': df.loc[(df['Churn']==churn), 'MonthlyCharges'],
            'y': df.loc[(df['Churn']==churn), 'tenure'],
            'name': churn, 'mode': 'markers',
        } for churn in churns
    ],
    'layout': {
        'title': 'Tenure vs Monthly Charges',
        'xaxis': {'title': 'Monthly Charges'},
        'yaxis': {'title': 'Tenure'}
    }
}
py.offline.iplot(fig)
```

```
In [15]: figs = []
for churn in churns:
    figs.append(
        go.Box(
            y = df.loc[(df['Churn']==churn), 'tenure'],
            name = churn
        )
    )
layout = go.Layout(
    title = 'Tenure',
    xaxis = {'title': 'Churn?'},
    yaxis = {'title': 'Tenure'},
    width=800,
    height=500
)
fig = go.Figure(data=figs, layout=layout)
py.offline.iplot(fig)
```

```
In [16]: figs = []
for churn in churns:
    figs.append(
        go.Box(
            y = df.loc[(df['Churn']==churn), 'MonthlyCharges'],
            name = churn
        )
    )
layout = go.Layout(
    title = 'MonthlyCharges',
    xaxis = {'title': 'Churn?'},
    yaxis = {'title': 'MonthlyCharges'},
    width=800,
    height=500
)
fig = go.Figure(data=figs, layout=layout)
py.offline.iplot(fig)
```

```
In [17]: _ = df.groupby('Churn').size().reset_index()
# sort values by 'tenure', ascending=True

data = [go.Bar(
    x = ['Churn'].tolist(),
    y = _[0].tolist(),
    marker=dict(
        color=['rgb(255,190,134,1)', 'rgb(142,186,217,1)']
    )
)]
layout = go.Layout(
    title = 'Churn Distribution',
    xaxis = {'title': 'Churn?'},
    width=800,
    height=500
)
fig = go.Figure(data=data, layout=layout)
py.offline.iplot(fig)
```

```
In [18]: df['class'] = df['Churn'].apply(lambda x: 1 if x == "Yes" else 0)
# features will be saved as X and our target will be saved as y
X = df[['tenure', 'MonthlyCharges']].copy()
X2 = df[['tenure', 'MonthlyCharges']].copy()
y = df['class'].copy()
```

```

In [ ]: start_time = time.time()

num_iter = 100000

intercept = np.ones((X.shape[0], 1))
X = np.concatenate((intercept, X), axis=1)
theta = np.zeros(X.shape[1])

for i in range(num_iter):
    h = sigmoid(X, theta)
    gradient = gradient_descent(X, h, y)
    theta = update_weight_loss(theta, 0.1, gradient)

print("Training time (Log Reg using Gradient descent):" + str(time.time() - start_time) + " seconds")
print("Learning rate: {} \n Iteration: {}".format(0.1, num_iter))

Training time (Log Reg using Gradient descent): 70.8485119342804 seconds
Learning rate: 0.1
Iteration: 100000

In [ ]: result = sigmoid(X, theta)

In [ ]: f = pd.DataFrame(np.around(result, decimals=6)).join(y)
f['pred'] = f[0].apply(lambda x: 0 if x < 0.5 else 1)
print("Accuracy (Loss minimization):")
f.loc[f['pred']==f['class']].shape[0] / f.shape[0] * 100

Accuracy (Loss minimization):
Out[ ]: 53.301150078091716

In [ ]: start_time = time.time()
num_iter = 100000

intercept2 = np.ones((X2.shape[0], 1))
X2 = np.concatenate((intercept2, X2), axis=1)
theta2 = np.zeros(X2.shape[1])

for i in range(num_iter):
    h2 = sigmoid(X2, theta2)
    gradient2 = gradient_descent(X2, h2, y) #np.dot(X.T, (h - y)) / y.size
    theta2 = update_weight_mle(theta2, 0.1, gradient2)

print("Training time (Log Reg using MLE):" + str(time.time() - start_time) + "seconds")
print("Learning rate: {} \n Iteration: {}".format(0.1, num_iter))

<ipython-input-2-2e0ea9337b29>:3: RuntimeWarning:
overflow encountered in exp

Training time (Log Reg using MLE): 81.35162234396335 seconds
Learning rate: 0.1
Iteration: 100000

In [ ]: result2 = sigmoid(X2, theta2)

<ipython-input-2-2e0ea9337b29>:3: RuntimeWarning:
overflow encountered in exp

In [ ]: from sklearn.linear_model import LogisticRegression

clf = LogisticRegression(fit_intercept=True, max_iter=100000)
clf.fit(df[['tenure', 'MonthlyCharges']], y)
print("Training time (sklearn's LogisticRegression module):" + str(time.time() - start_time) + " seconds")
print("Learning rate: {} \n Iteration: {}".format(0.1, num_iter))

Training time (sklearn's LogisticRegression module): 83.02515387535095 seconds
Learning rate: 0.1
Iteration: 100000

In [ ]: result3 = clf.predict(df[['tenure', 'MonthlyCharges']])

In [ ]: print("Accuracy (sklearn's Logistic Regression):")
f3 = pd.DataFrame(result3).join(y)
f3.loc[f3[0]==f3['class']].shape[0] / f3.shape[0] * 100

Accuracy (sklearn's Logistic Regression):
Out[ ]: 78.44668465142695

```


Lab5

Date: 24th May, 2024

Title: Build Support vector machine model for a given dataset

Algorithm (Handwritten)

24/5/24 LAB5

- 1) Support Vector Machine
- 2) Define kernel function
Eg $K(x_1, x_2) = x_1 \cdot x_2$
- 3) Solve the quadratic programming problem (QP) to find the x
- 4) Compute the bias
- 5) Identify the support vectors
- 6) Make prediction

OUTPUT

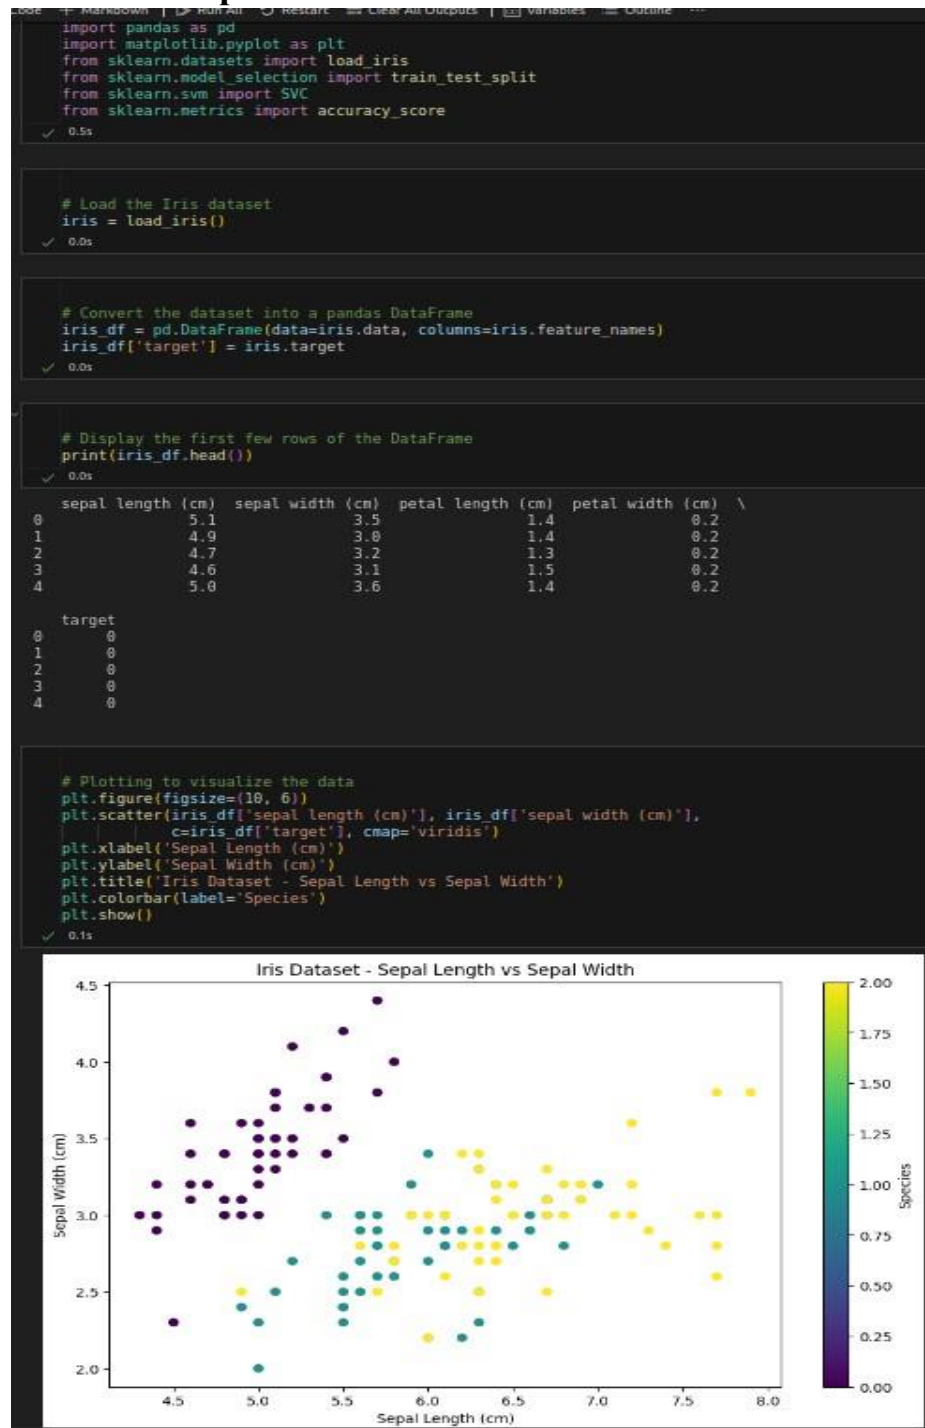
```
Model = svm()
Model.fit(X_train, Y_train)
prediction = model.predict(X_test)
accuracy = (y_test, prediction)
```

0.98230086

→ Model predict ([-0.47096, -0.1604584, -0.47181 --
-0.244122, -0.19956318, 0.1832044,
-0.19695794])

array(0)

Code and output



```

# Splitting the dataset into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(iris.data, iris.target, test_size=0.3, random_state=42)

# Creating and training the SVM classifier
svm_classifier = SVC(kernel='linear')
svm_classifier.fit(X_train, y_train)

# Predicting the labels for the test set
y_pred = svm_classifier.predict(X_test)

# Calculating the accuracy of the model
accuracy = accuracy_score(y_test, y_pred)
print("Accuracy of SVM Classifier:", accuracy)

Accuracy of SVM Classifier: 1.0

y_pred
array([1, 0, 2, 1, 1, 0, 1, 2, 1, 1, 2, 0, 0, 0, 0, 1, 2, 1, 1, 2, 0, 2,
       0, 2, 2, 2, 2, 0, 0, 0, 0, 1, 0, 0, 2, 1, 0, 0, 0, 2, 1, 1, 0,
       0])

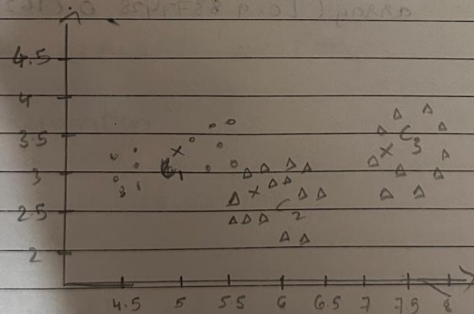
```

Title: Build K-Means algorithm to cluster a set of data stored in a .CSV file

Algorithm

- 2) K means clustering Algorithm
- 1) Select number K decide the no. of elements.
- 2) Select random k points or centroids.
- 3) Assign each point to the nearest centroid, which will form the predefined cluster.
- 4) Calculate the variance and place a new centroid of each cluster.
- 5) Repeat step 3, reassign the centroid.
- 6) If any rearrangement occurs perform step 4 else finish.
- 7) The model is ready.

OUTPUT:



Code and output

```
import matplotlib.pyplot as plt
from sklearn import datasets
from sklearn.cluster import KMeans
import pandas as pd
import numpy as np
```

```
# import some data to play with
iris = datasets.load_iris()
X = pd.DataFrame(iris.data)
X.columns = ['Sepal_Length', 'Sepal_Width', 'Petal_Length', 'Petal_Width']
y = pd.DataFrame(iris.target)
y.columns = ['Targets']
```

```
# Build the K Means Model
model = KMeans(n_clusters=3)
model.fit(X) # model.labels_ : Gives cluster no for which samples belongs
```

```
/usr/local/lib/python3.10/dist-packages/sklearn/cluster/_kmeans.py:878: FutureWarning:
warnings.warn(
```


```
+ KMeans
KMeans(n_clusters=3)
```

```
# Visualise the clustering results
plt.figure(figsize=(14,14))
colormap = np.array(['red', 'lime', 'black'])
```

<Figure size 1400x1400 with 0 Axes>


```
# Plot the Original Classifications using Petal features
plt.subplot(2, 2, 1)
plt.scatter(X.Petal_Length, X.Petal_Width, c=colormap[y.Targets], s=40)
plt.title('Real Clusters')
plt.xlabel('Petal Length')
plt.ylabel('Petal Width')
```

Text(0, 0.5, 'Petal Width')



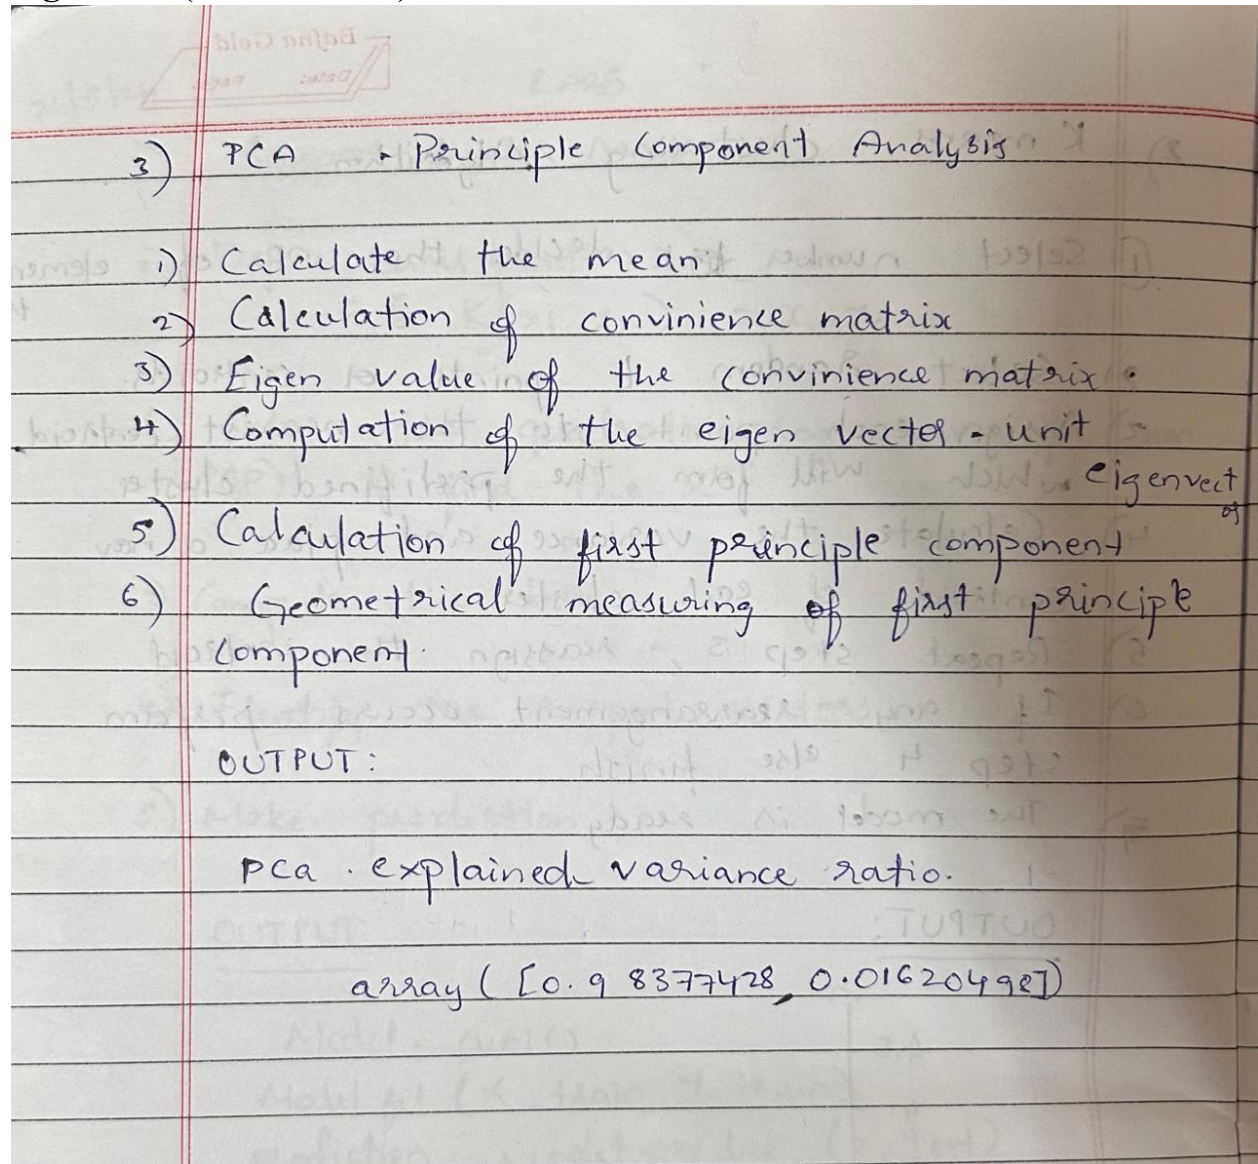
```
# Plot the Models Classifications
plt.subplot(2, 2, 2)
plt.scatter(X.Petal_Length, X.Petal_Width, c=colormap[model.labels_], s=40)
plt.title('K-Means Clustering')
plt.xlabel('Petal Length')
plt.ylabel('Petal Width')
```

Text(0, 0.5, 'Petal Width')



Title: Implement Dimensionality reduction using Principle Component Analysis (PCA)
Method

Algorithm (Handwritten)



Code and output

```
import matplotlib.pyplot as plt
from sklearn import datasets
from sklearn.decomposition import PCA
from sklearn.preprocessing import StandardScaler
import pandas as pd
import numpy as np

# Load the iris dataset
iris = datasets.load_iris()
X = pd.DataFrame(iris.data, columns=['Sepal.Length', 'Sepal.Width', 'Petal.Length', 'Petal.Width'])
y = pd.DataFrame(iris.target, columns=['Targets'])

# Standardize the data
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)


# Apply PCA
pca = PCA(n_components=2)
X_pca = pca.fit_transform(X_scaled)

# Convert PCA result to a DataFrame
X_pca_df = pd.DataFrame(X_pca, columns=['PCA1', 'PCA2'])

# Add the target column for visualization
X_pca_df['Targets'] = y.Targets

# Visualize the PCA result
plt.figure(figsize=(10, 7))
colormap = np.array(['red', 'lime', 'black'])

# Plot the PCA transformed data
plt.scatter(X_pca_df.PCA1, X_pca_df.PCA2, c=colormap[X_pca_df.Targets], s=40)
plt.title('PCA of Iris Dataset')
plt.xlabel('Principal Component 1')
plt.ylabel('Principal Component 2')
plt.show()
```



Lab6

Date: 31st May, 2024

Title: Build Artificial Neural Network model with back propagation on a given dataset.

Refer: <https://docs.google.com/presentation/d/11UE61G27eOAYnhc8ctHAqoEaeYLrVhoT/edit?usp=sharing&ouid=117926028109390959744&rtpof=true&sd=true>

Algorithm (Handwritten)

31/05/24 LAB-6 Bafna Gold Date: Page:

Build an artificial ~~neural~~ neural network model with back propagation.

Algorithm:

- * Initialize parameters
- * Normalize input features
- * Normalize the output
- * Set hyper parameters: no. of epochs, no. of neurons
- * Define activation functions
- Training the network
 - Forward propagation
 - * Compute input to hidden layer
 - * Add bias
 - * apply activation function
 - Backpropagation
 - * Compute error
 - * compute gradient
 - * Compute delta
- Update weights and biases

Code and output

```
import numpy as np
x = np.array([[2,9],[1,5],[3,6]],dtype = float)
y = np.array([92],[86],[89]),dtype = float)
x = x/np.amax(x,axis=0)
y = y/100

#Variable Initialization
epoch = 5000
lr = 0.1
inputlayer_neurons = 2
hiddenlayer_neurons = 3
output_neurons = 1

# weight and bias Initialization
wh = np.random.uniform(size=(inputlayer_neurons,hiddenlayer_neurons))
bh = np.random.uniform(size=(1,hiddenlayer_neurons))
wout = np.random.uniform(size=(hiddenlayer_neurons,output_neurons))
bout = np.random.uniform(size=(1,output_neurons))

#sigmoid function
def sigmoid(x):
    return 1/(1+np.exp(-x))

# Derivative of Sigmoid
def der_sigmoid(x):
    return x*(1-x)

# Draws a random range of numbers uniformly of dim x*y
for i in range(epoch):

    # forward propagation
    hinpl = np.dot(x,wh)
    hinp = hinpl + bh
    hlayer_act = sigmoid(hinp)
    outinp1 = np.dot(hlayer_act,wout)
    outinp = outinp1 + bout
    output = sigmoid(outinp)

    # Backpropagation
    E0 = y - output
    outgrad = der_sigmoid(output)
    d_output = E0*outgrad
    EH = d_output.dot(wout.T)
```

```
# how much hidden layer weights contributed to error
hiddengrad = der_sigmoid(hlayer_act)
d_hiddenlayer = EH*hiddengrad

#dotproduct of nextlayererror and current layer op
wout += hlayer_act.T.dot(d_output)*lr
wh += x.T.dot(d_hiddenlayer)*lr

print("Input: \n" + str(x))
print("Actual output: \n" + str(y))
print("Predicted Output: \n",output)
```

```
Input:
[[0.66666667 1.         ]
 [0.33333333 0.55555556]
 [1.         0.66666667]]
Actual output:
[[0.92]
 [0.86]
 [0.89]]
Predicted Output:
[[0.80056875]
 [0.79393831]
 [0.80112347]]
```

Title: Implement Random forest ensemble method on a given dataset.

Ref- <https://towardsdatascience.com/random-forest-in-python-24d0893d51c0>

Algorithm (Handwritten)

Implement Random Forest Ensemble Method

Algorithm:

- 1) Import necessary libraries
- 2) Load & insert data
- 3) Pre process the data as in separating features and strengths
- 4) Split the data to training and test samples. Use 0.4 to allocate 40% of data to testing and use rest for training.
- 5) Initialize random forest classifier and train it using fit method.
- 6) Make predictions on test sample using method predict.
- 7) Evaluate the model

OUTPUT:

Accuracy: 0.98

Confusion Matrix:

$\begin{bmatrix} 23 & 0 & 0 \end{bmatrix}$

$\begin{bmatrix} 0 & 19 & 0 \end{bmatrix}$

$\begin{bmatrix} 0 & 1 & 17 \end{bmatrix}$

Code and output

```
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score, classification_report, confusion_matrix
from sklearn import datasets

# Load the data
iris_data = datasets.load_iris()

X = pd.DataFrame(iris_data.data, columns=['Sepal_Length', 'Sepal_Width', 'Petal_Length', 'Petal_Width'])
y = pd.DataFrame(iris_data.target, columns=['Targets'])

# Check the info of the modified data
# print(iris_data.info())

# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.4, random_state=42)

# Initialize the Random Forest classifier
rf_classifier = RandomForestClassifier(n_estimators=100, random_state=42)

# Fit the classifier to the training data
rf_classifier.fit(X_train, y_train)

# Predict on the test data
y_pred = rf_classifier.predict(X_test)

# Evaluate the classifier
accuracy = accuracy_score(y_test, y_pred)
print(f"Accuracy: {accuracy:.2f}")

# Print classification report
print("Classification Report:")
print(classification_report(y_test, y_pred))

# Print confusion matrix
print("Confusion Matrix:")
print(confusion_matrix(y_test, y_pred))
```

✓ 0.7s

Accuracy: 0.98

Classification Report:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	23
1	0.95	1.00	0.97	19
2	1.00	0.94	0.97	18
accuracy			0.98	60
macro avg	0.98	0.98	0.98	60
weighted avg	0.98	0.98	0.98	60

Confusion Matrix:

```
[[23  0  0]
 [ 0 19  0]
 [ 0  1 17]]
```


Title: Implement Boosting ensemble method on a given dataset

Algorithm (Handwritten)

Implement Boosting ensemble method on a given dataset.

Algorithm:

- 1) Import Libraries
- 2) Load the dataset
- 3) Data pre processing involves separation of features and dataset
- 4) Split the dataset to train and test samples
- 5) Initialize the adaboost classifier with specified no. of estimates and base estimators
- 6) Train the model using the training data.
- 7) Make predictions for test sample using trained model.
- 8) Evaluate the model

OUTPUT:

Metrics. accuracy score:

0.983333333333

Code and output

```
[10] from sklearn.linear_model import LogisticRegression
     from sklearn.ensemble import AdaBoostClassifier
     from sklearn import metrics
     from sklearn import datasets

[11] import pandas as pd
     import matplotlib.pyplot as plt
     from sklearn.model_selection import train_test_split

[12] # Load the Iris dataset
     iris = datasets.load_iris()
     X = pd.DataFrame(iris.data, columns=['Sepal_Length', 'Sepal_Width', 'Petal_Length', 'Petal_Width'])
     y = pd.DataFrame(iris.target, columns=['Targets'])

[13] X_train,X_test,y_train,y_test=train_test_split(X,y,test_size=0.4,random_state=42)

[14] mylogregmodel = LogisticRegression()

[15] adabc = AdaBoostClassifier(n_estimators = 150, estimator = mylogregmodel, learning_rate = 1)

[16] model = adabc.fit(X_train, y_train)
     /usr/local/lib/python3.10/dist-packages/sklearn/utils/validation.py:1143: DataConversionWarning: A
     y = column_or_1d(y, warn=True)

[17] y_pred = model.predict(X_test)

[18] metrics.accuracy_score(y_test, y_pred)

     0.9833333333333333
```